




LLMs Compress World Knowledge




What is RAG?

Augmenting LLM with specialized and mutable knowledge base



How it works?


2 key elements: retriever and knowledge base




How it works?

2 key elements: **retriever** and knowledge base

Text Embeddings



Text Embedding-based Search



How it works?

2 key elements: retriever and **knowledge base**

